

CSE 421/521 - Operating Systems Spring 2018

LECTURE - XVI

PROJECT-2 DISCUSSION

Tevfik Koşar

University at Buffalo
April 3rd, 2018

Pintos Projects

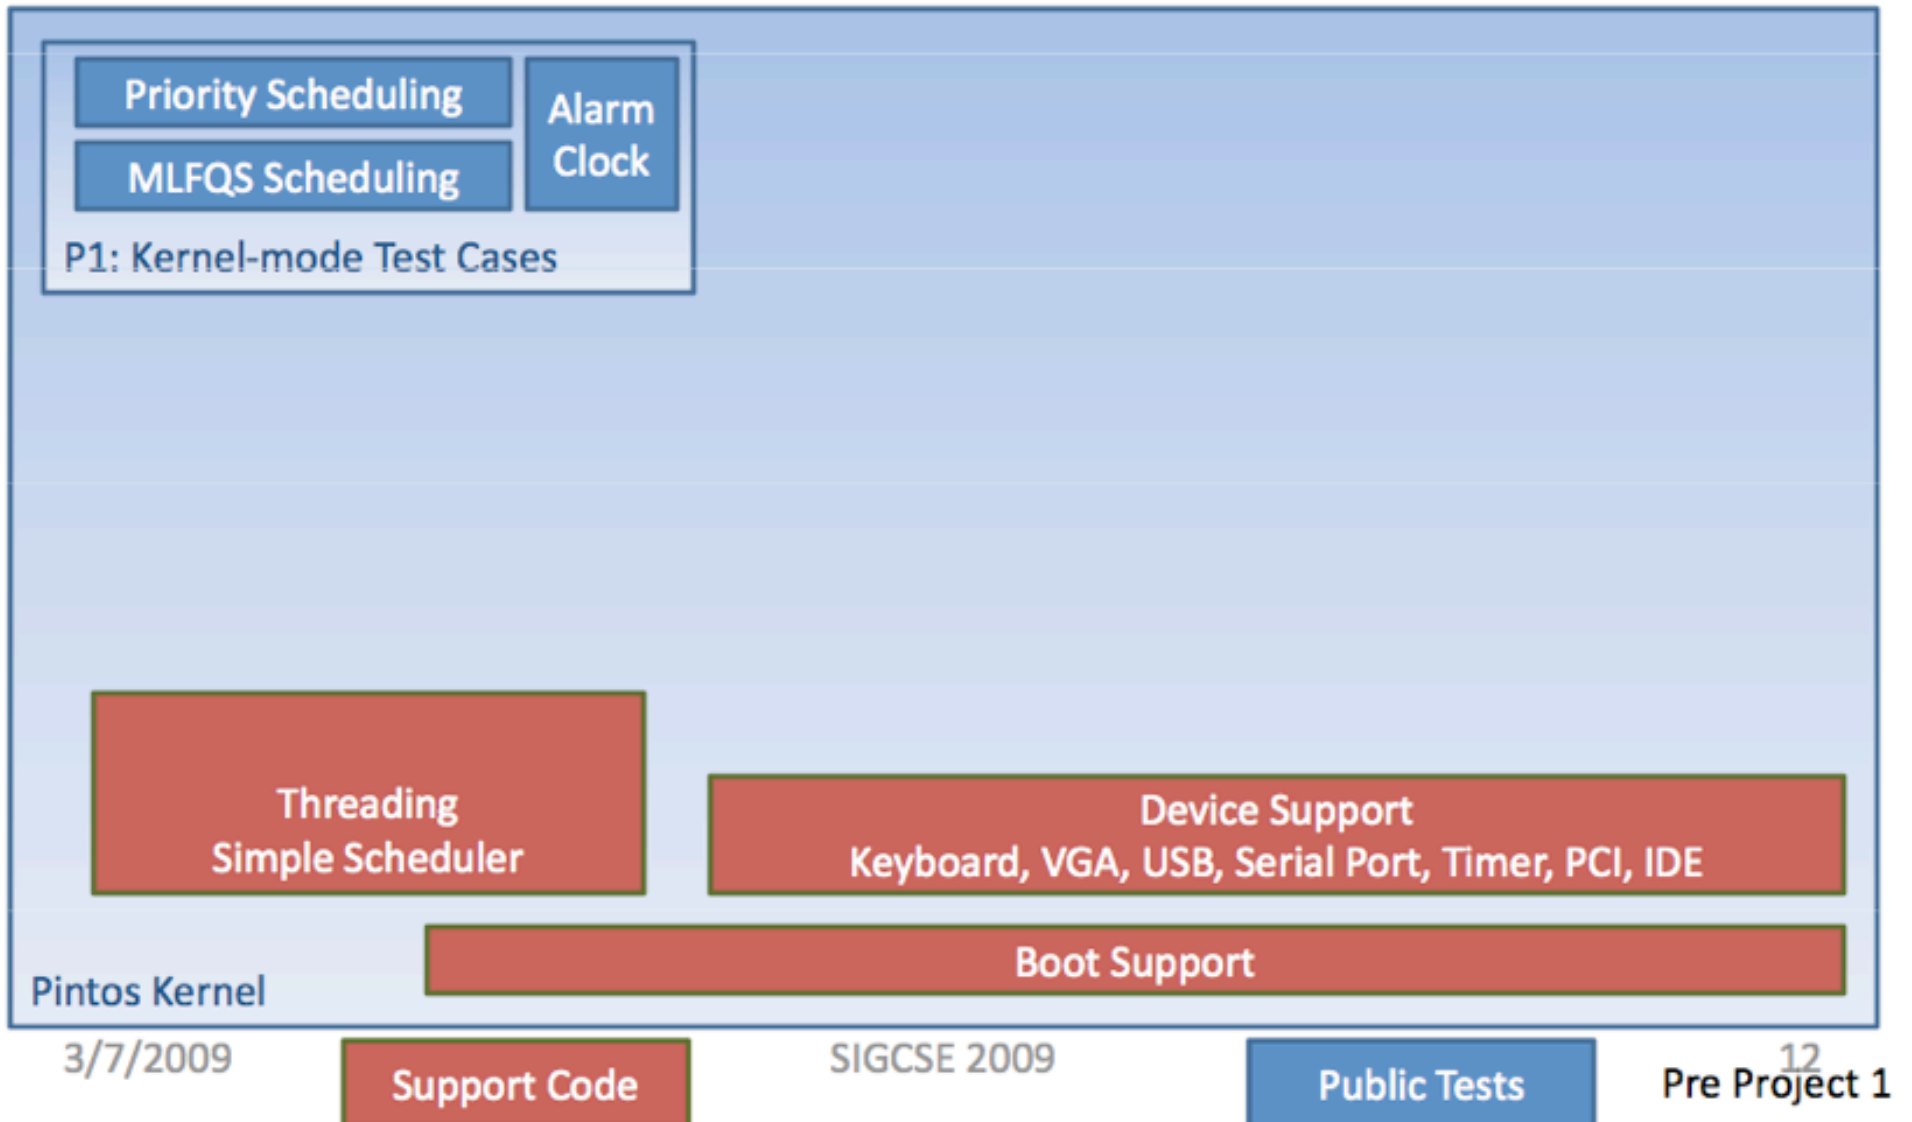
1. Threads

2. User Programs <-- CSE 421/521 Project 2

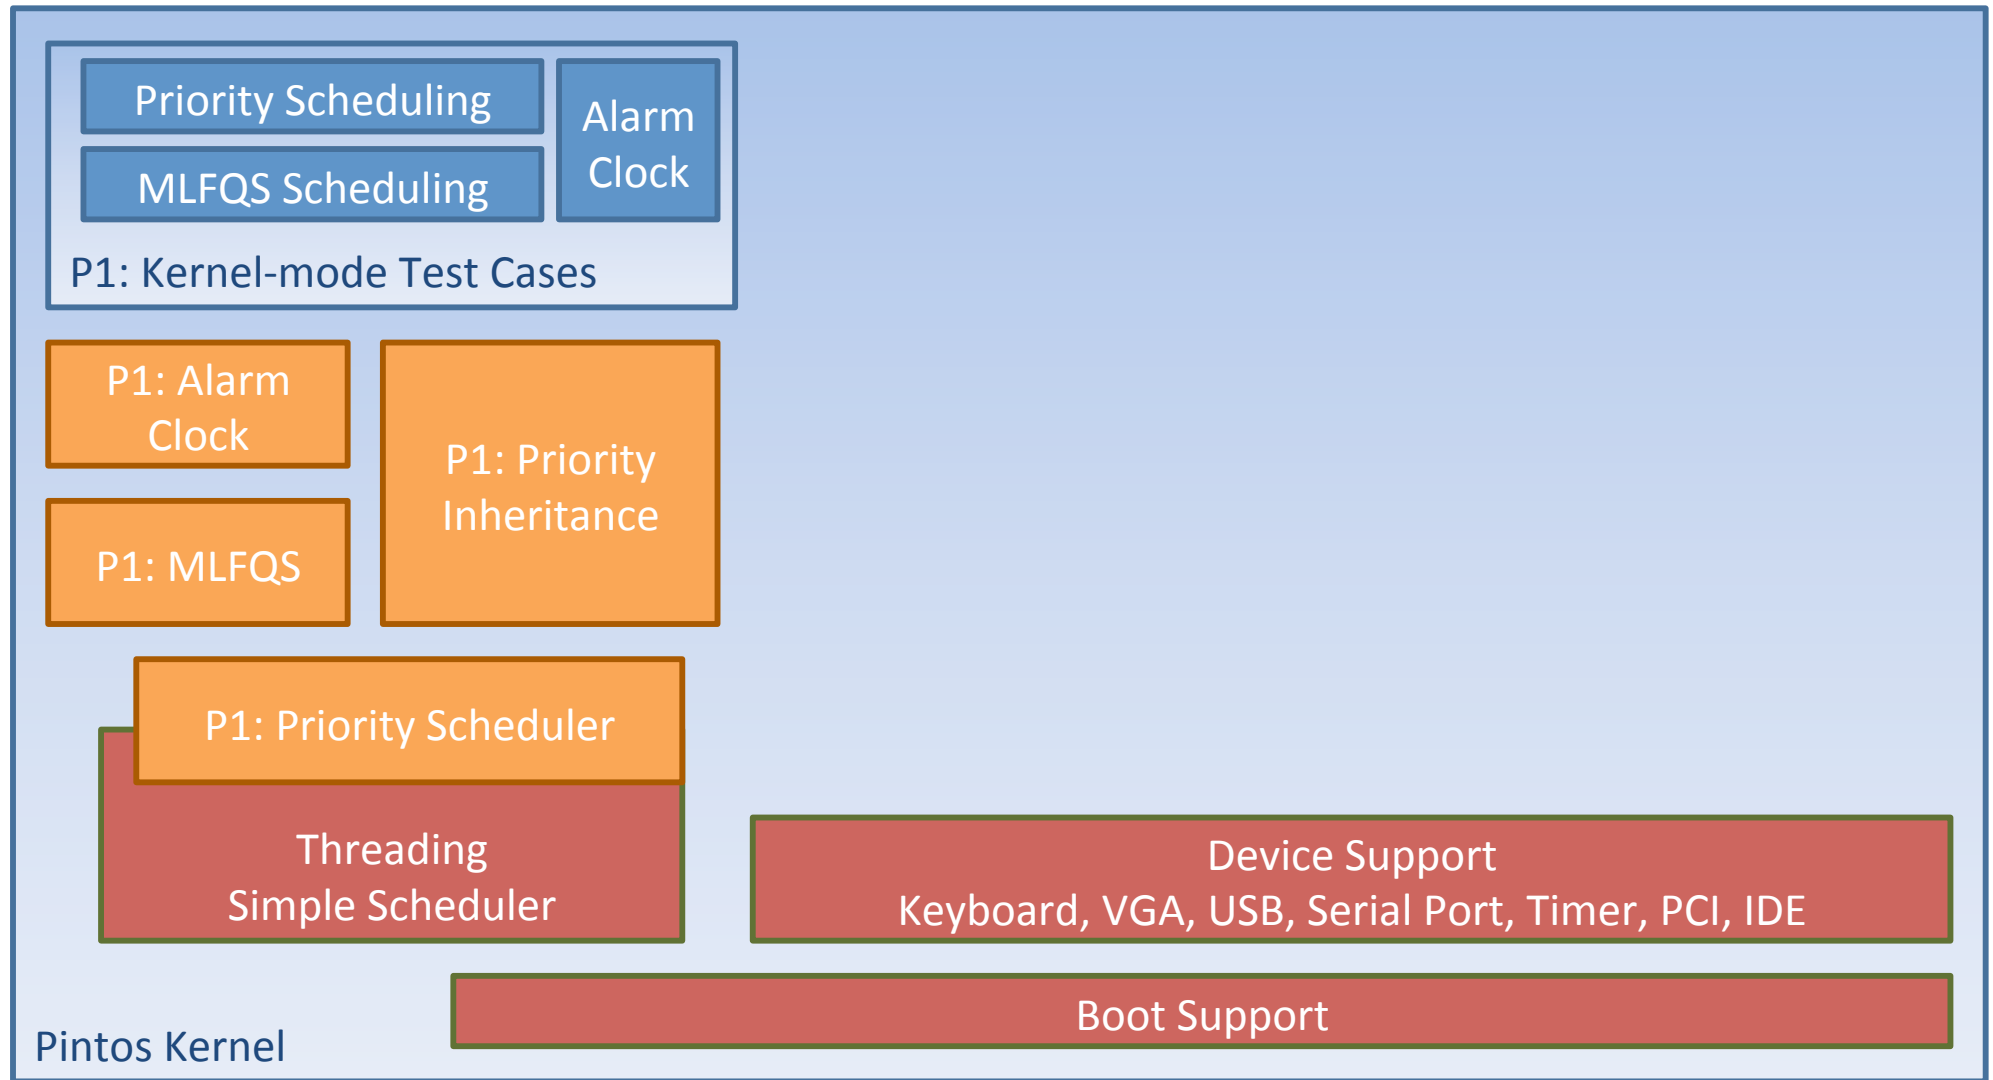
3. Virtual Memory

4. File Systems

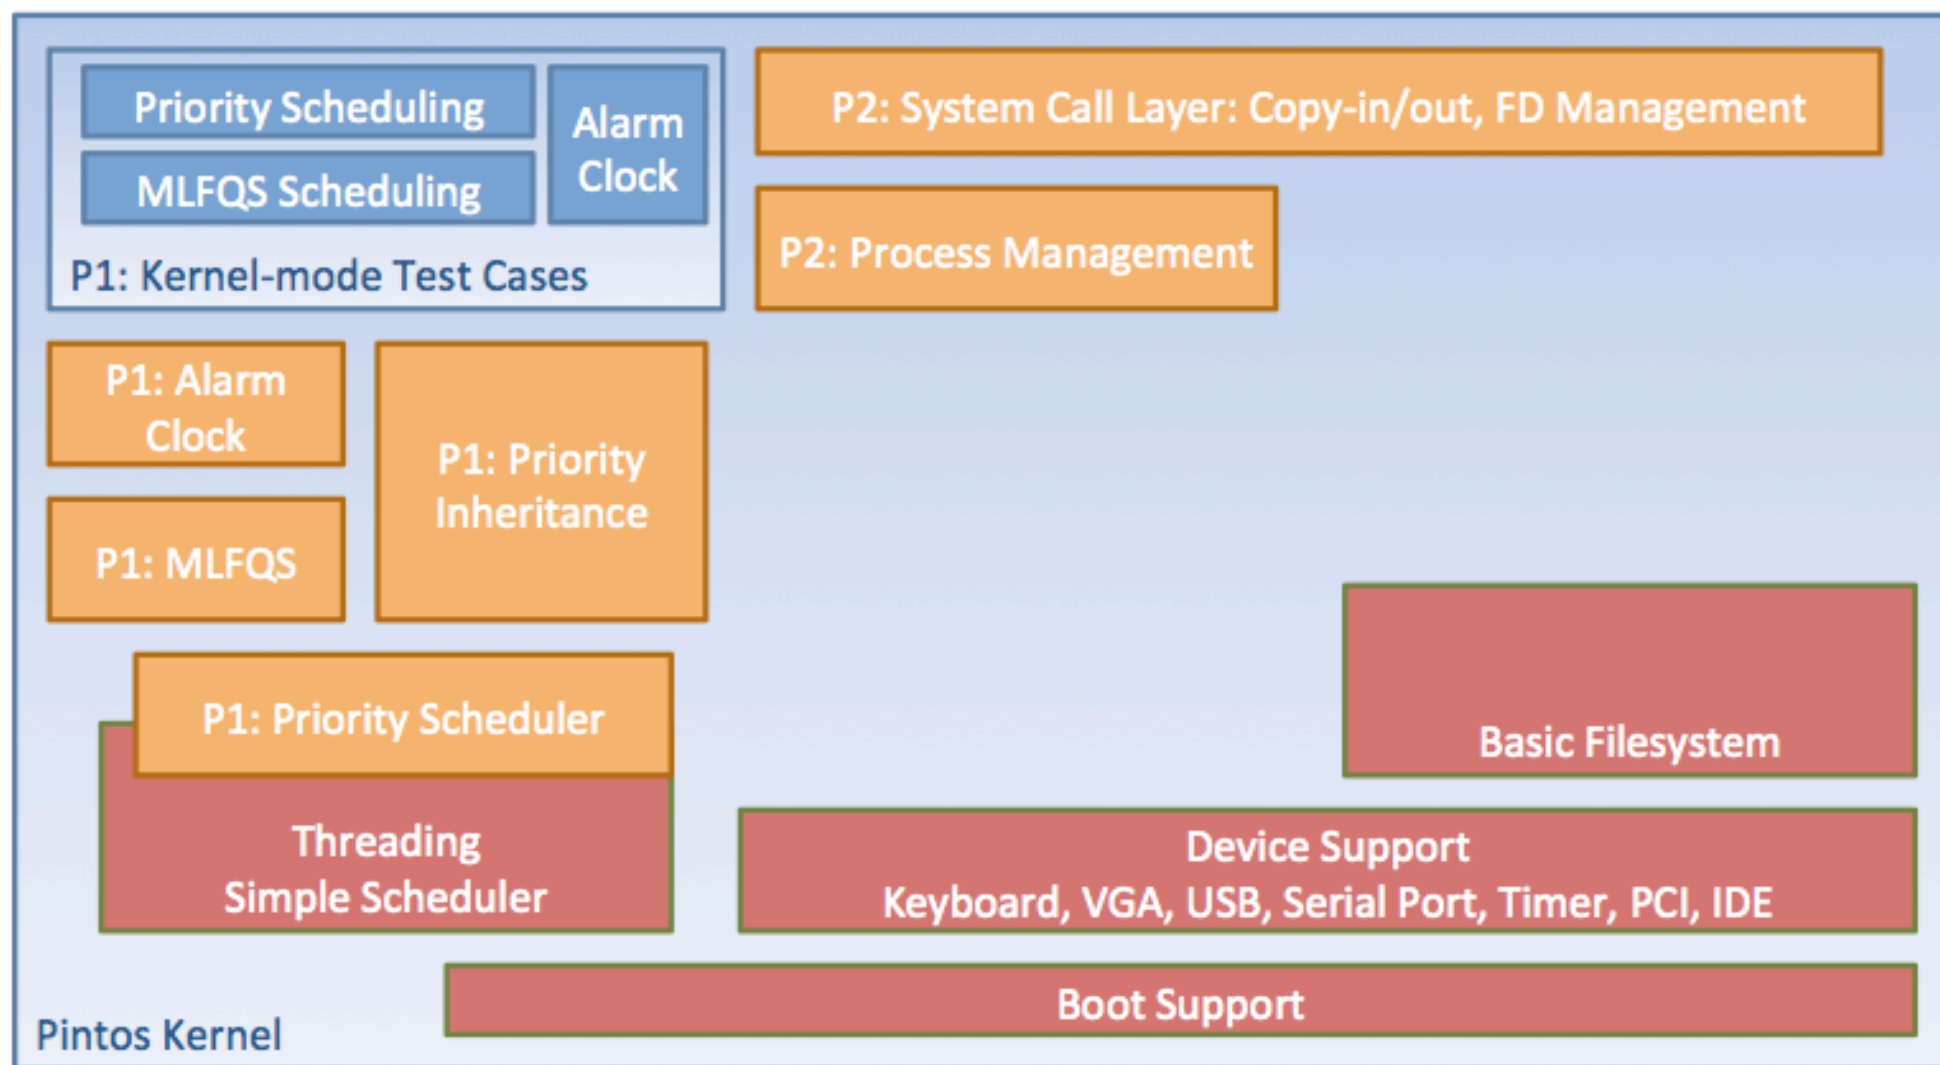
Yo were provided with this (pre prj-1)



You implemented this (post prj-1)



You will now implement this (post prj-2)



Overview (1)

- The existing Pintos code already supports loading and running user programs, but no argument passing and no interactivity (I/O) is possible.
- In this project, you will enable programs to support arguments and interact with the OS via system calls.

Overview (2)

- This project requires an understanding of:
 - How user programs run in general
 - Distinctions between user and kernel virtual memory
 - The system call infrastructure / file system interface

Example (1)

- Sample C program without arguments and systems calls:

```
main()
{
    int n=10, fact=1;
    for (int i=1; i<=10; i++)
    {
        fact = fact * i;
    }
}
```


Example (2)

- In C, a user program test.c can pass argument

- ```
int main(int argc, char* argv[])
{
 for(int i=0; i<argc; i++)
 {
 char* arg = argv[i];
 }
}
```

- ```
./test arg1 arg2 ...
```

Example (3)

- test.c can call system libraries

- #include <stdio.h>

- int main()

- {

- FILE* p_file = fopen("myfile.txt","w");

- if (p_file != NULL) fputs("fopen", p_file);

- fclose(p_file);

- }

- Get fopen, fputs, fclose by system calls

- Pintos need you to implement

- Argument passing

- System calls

Example (4)

- What happens when a user wants to run the following program in the Unix shell?

```
myth13:~$ cp -r /usr/bin/temp .
```

- 1) Shell parses user input
 - 2) Shell calls `fork()` and `execve("cp", argv, env)`
 - 3) `cp` uses file system interface to copy files
 - 4) `cp` may print messages to `stdout`
 - 5) `cp` exits
- These interactions require system calls

Project Goals

You will need to implement the following:

- Passing command-line arguments to programs
- Safe memory access
- A set of system calls
 - Long list, in section 3.3.4 of the Pintos docs
- Process termination messages
- Denying writes to files in use as executables

File System Access in Pintos (1)

- A basic file systems is already provided in Pintos
- You will need to interact with this file system
- But, **do not modify the file system!**
- UNIX like semantics: open, close, read, write...
- Files to take a look at: '*filesys.h*' & '*file.h*'

File System Access in Pintos (2)

Pintos File system limitations:

- No internal synchronization
- File size is fixed at creation time
- File data is allocated as a single extent (i.e., in a contiguous range of sectors on disk)
- No subdirectories
- File names are limited to 14 characters
- A system crash may corrupt the disk

File System Access in Pintos (2)

■ Using the Pintos file system

```
$ pintos-mkdisk fs.dsk 2
```

– Creates a 2MB disk named "fs.dsk"

```
$ pintos -f -q
```

– Formats the disk (-f) and exits as soon as the format is done (-q)

```
$ pintos -p ../../examples/echo -a echo -- -q
```

– Put the file "../../examples/echo" to the Pintos file system under the name "echo"

```
$ pintos -q run 'echo x'
```

– Run the executable file "echo", passing argument "x"

```
$ pintos -fs-disk=2 -p ../../examples/echo -a  
echo -- -f -q run 'echo x'
```

Step 1: Preparation

READ:

- Sections 2.3 - 2.4 from Silberschatz
 - Systems Calls
- Chapter 8-9 from Silberschatz
 - Memory & Virtual Memory Management
- Lecture slides on Memory & Virtual Memory Management
 - Lectures 14 & 15
- From Pintos Documentation:
 - Section 3: User Programs

Step 2: Setting Up Pintos (1)

- **Project-2 does not depend on project-1**

No code from project-1 is required for this assignment

You can use your existing code base if you wish,

Or, if you think you have totally messed up with project-1, you can grab a fresh copy of the Pintos code base, following the instructions on the project description.

Step 2: Setting Up Pintos (2)

- Use the Pintos **VM** we have prepared for you:

<http://ftp.cse.buffalo.edu/CSE421/UB-pintos.ova>

- It requires **Virtualbox** software

==> will work on most Linux, Windows, Mac systems

<https://www.virtualbox.org/wiki/Downloads>

- Detailed setup instructions are available on Piazza.

Step 3: Implementation

1. Argument Passing
2. User Memory Access
3. System Call Infrastructure
4. Implement 13 system calls (Start with Exit and Write)
5. Denying Writes to Executables

Task 0: Get Familiar with the Code

- The first task is to read and understand the code for the initial userprog system (under the “[pintos/src/userprog/](#)” directory).
- For a brief overview of the files in the “userprog/” directory, please see “Section 3.1.1. Source Files” in the Pintos Reference Guide
- You will be mostly working on:
 - /userprog/syscall.c
 - /userprog/process.c

Task 1: Argument Passing

- Before a user program starts executing, the kernel must push the function's arguments onto the stack.
- This involves breaking the command-line input into individual words.
- Consider `"/bin/ls -l foo bar" => "/bin/ls", "-l", "foo", "bar"`
- Implement the string parsing however you like in `process_execute()`
 - one solution is to use `strtok_r()` in `lib/string.c`

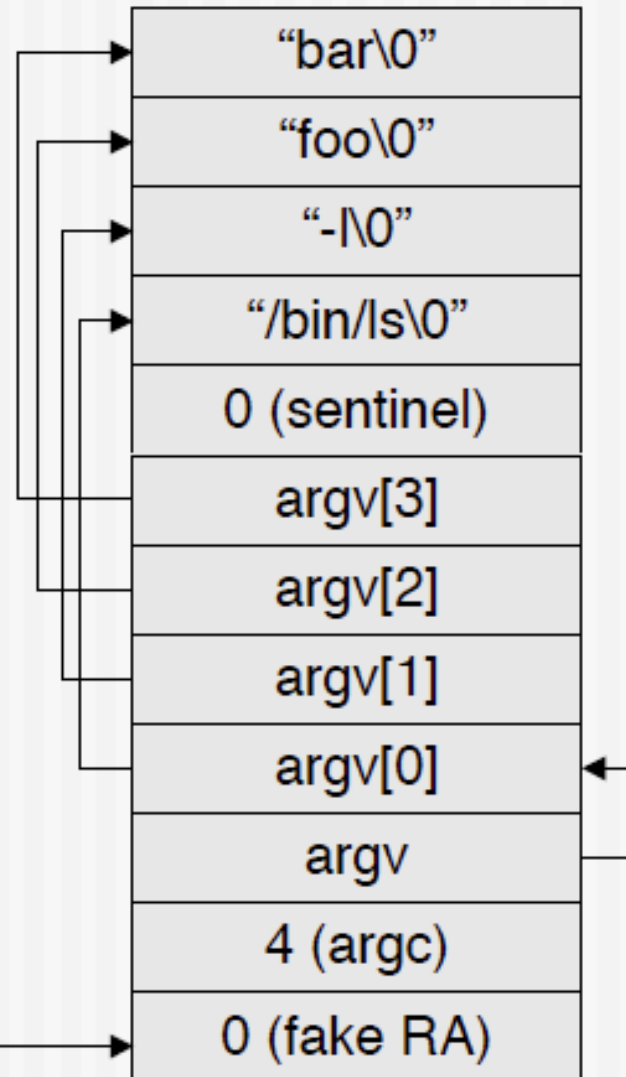
Argument Passing (cont.)

```
/bin/ls -l  
foo bar
```

argc = 4

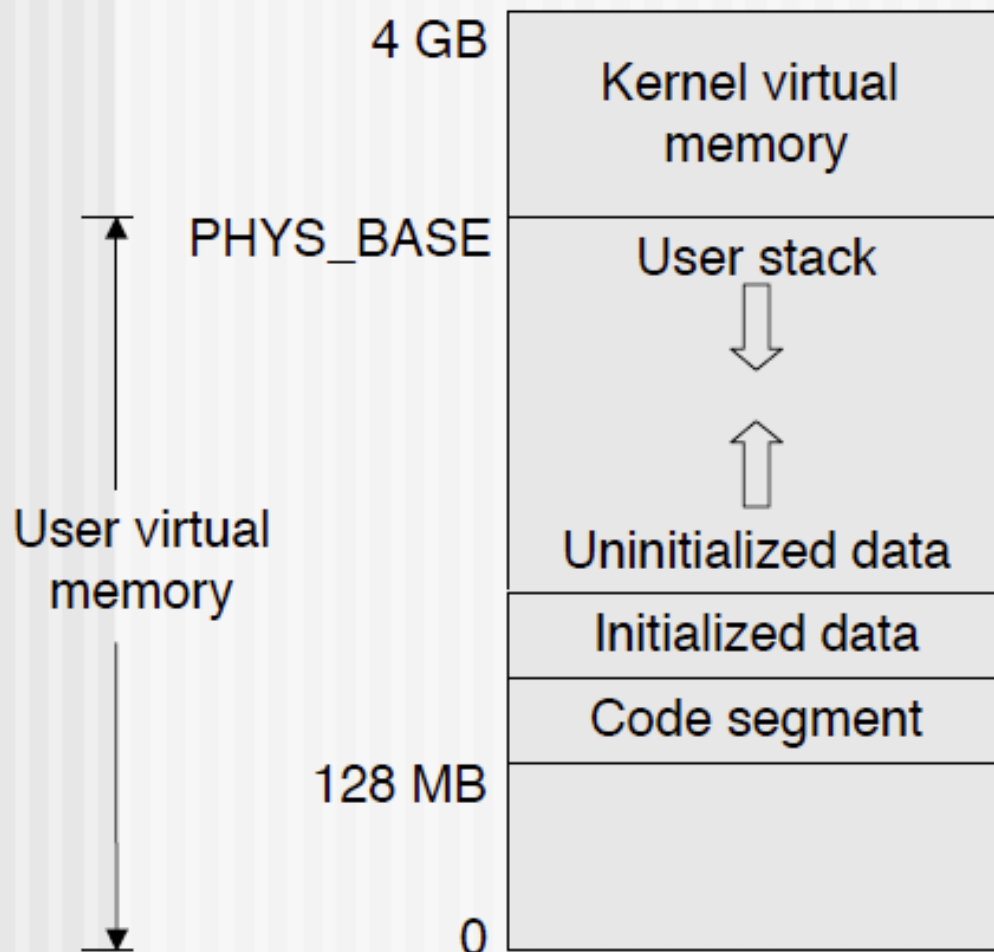
(Note: word
alignment not
shown)

stack pointer



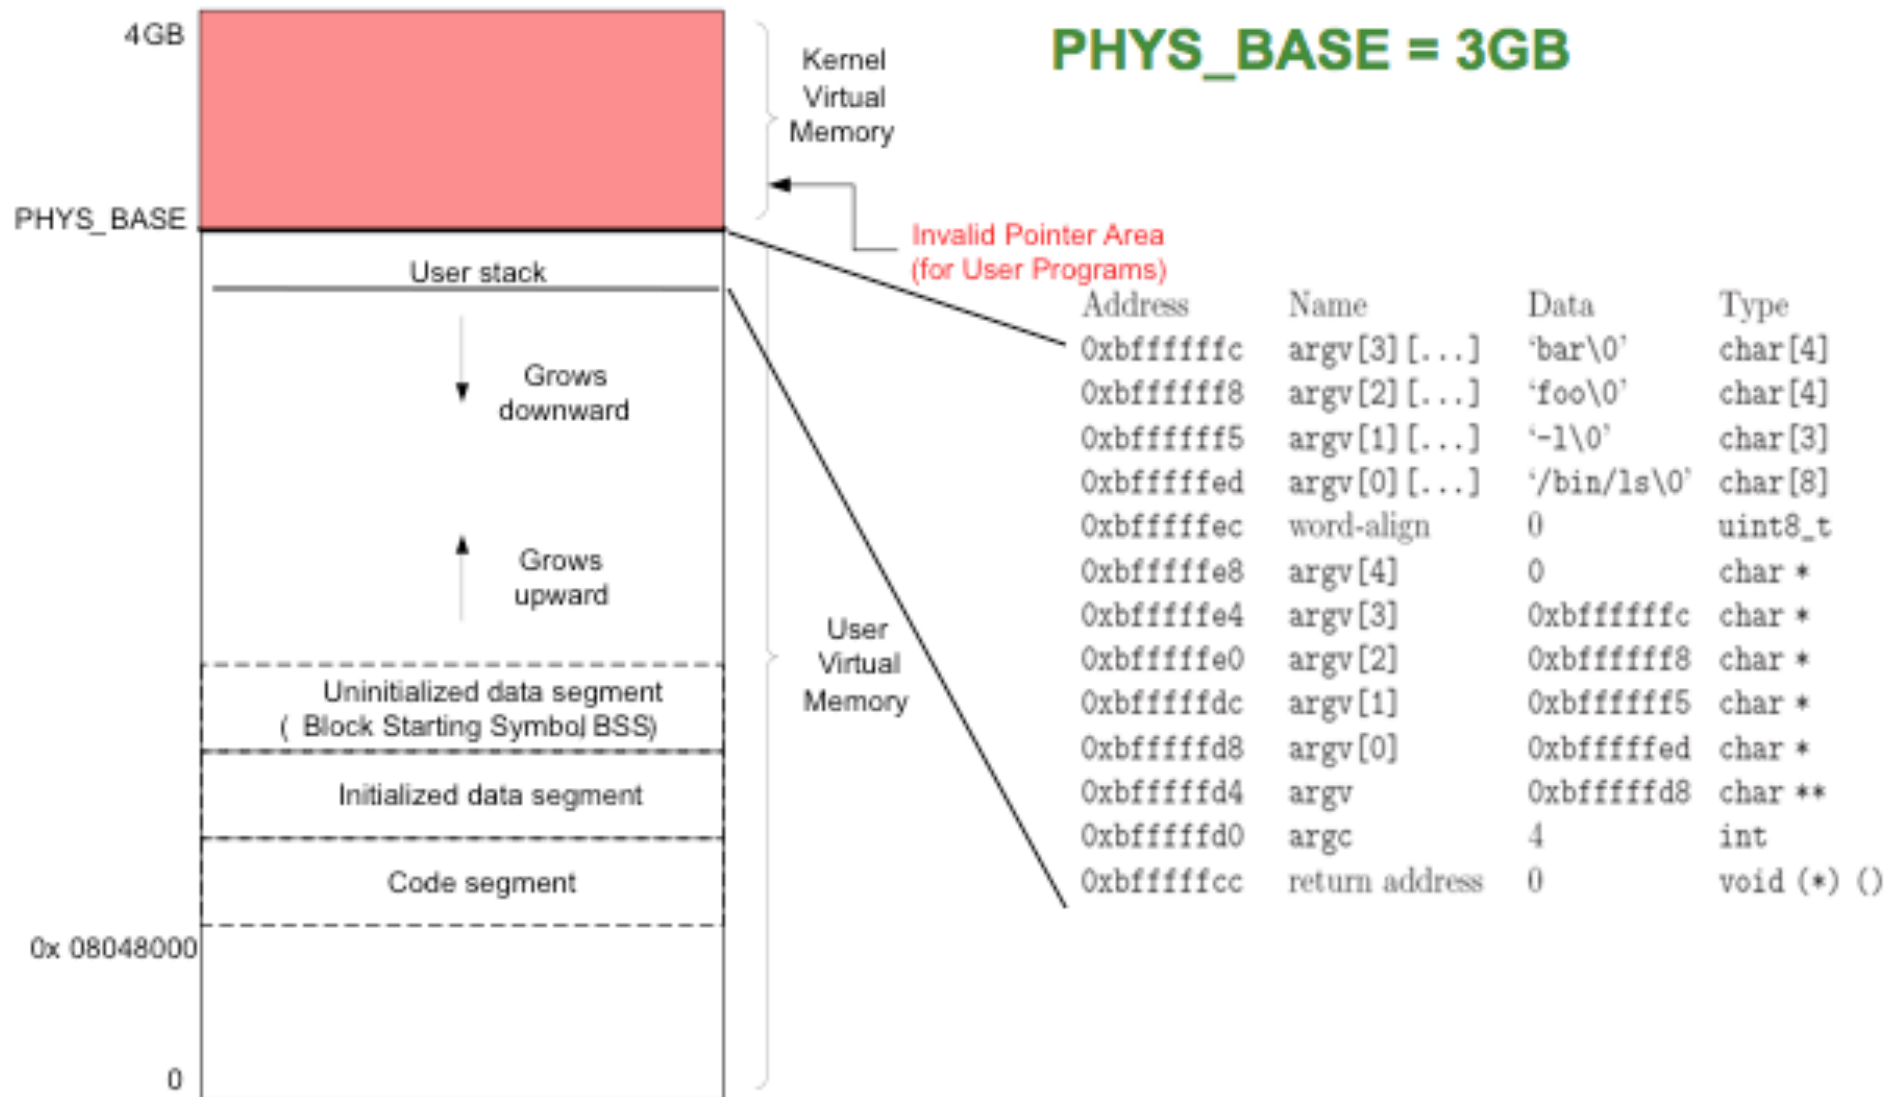
- Push the words onto the stack
- Push a null pointer sentinel
- Push the address of each word in right-to-left order
- Push argv and argc
- Push 0 as a fake return address
- hex_dump() function in <stdio.h> may be useful for seeing the layout of your stack

Virtual Memory Layout



- Virtual memory is divided between user and kernel virtual memory.
- Each user process has its own mapping of user virtual addresses (the page directory).
- User processes are not allowed to access kernel memory or unmapped user virtual addresses
 - This causes **page faults**
- The kernel can also page fault if it accesses an unmapped user virtual address

Memory Layout



Argument Passing (cont.)

PHYS_BASE = 0xc0000000	Kernel space				<div>command: /bin/ls -l foo bar</div>	
0xbfffffffcc	'b'	'a'	'r'	00		
0xbfffffff8	'f'	'o'	'o'	00		
0xbfffffff4	00	'_'	'l'	00		
padding → 0xbffffffc	'n'	'/'	'l'	's'		
0xbffffffe8	00	'/'	'b'	'i'	0	argv[4]
0xbffffffe4	fc	ff	ff	bf	0xbffffffc	argv[3]
0xbffffffe0	f8	ff	ff	bf	0xbfffffff8	argv[2]
0xbffffffdc	f5	ff	ff	bf	0xbffffff5	argv[1]
0xbffffffd8	ed	ff	ff	bf	0xbffffffed	argv[0]
0xbffffffd4	d8	ff	ff	bf	0xbffffffd8	argv
0xbffffffd0	04	00	00	00	4	argc
0xbffffffcc	return address					

Hints

- Pintos currently lacks argument passing!
- Change `*esp = PHYS_BASE` to `*esp = PHYS_BASE - 12` in `setup_stack()` to get started
- Change `process_execute()` in `process.c` to process multiple arguments
- Could limit the arguments to fit in a page(4 kb)
- String Parsing: `strtok_r()` in `lib/string.h`

```
pgm.c
main(int argc,
      char *argv[]) {
    ...
}

$ pintos run 'pgm alpha beta'
argc = 3
argv[0] = "pgm"
argv[1] = "alpha"
argv[2] = "beta"
```

User Programs in Pintos

`threads/init.c`

- `main()` => `run_actions(argv)` after booting
- `run_actions` => `run_task(argv)`
 - the task to run is `argv[1]`
- `run_task` => `process_wait(process_execute(task))`

`userprog/process.c`

- `process_execute` creates a thread that runs `start_process(filename...)` => `load(filename...)`
- `load` sets up the stack, data, and code, as well as the start address

Task 2: Safe Memory Access

- The kernel will often access memory through user-provided pointers
- These pointers can be problematic:
 - null pointers, pointers to unmapped user virtual memory, or pointers to kernel addresses
 - If a user process passes these to the kernel, you must kill the process and free its resources (locks, allocated memory)
- Be aware of potential problems with buffers, strings, and pointers

Hints

Two approaches to solving this problem:

- Verify every user pointer before dereference (simpler)
 - Is it in the user's address space, ie below PHYS_BASE? (look at `is_user_vaddr` in `threads/vaddr.h`)
 - Is it mapped or unmapped? (look at `pagedir_get_page()` in `userprog/pagedir.c`)
 - These checks apply to buffers as well
- Modify fault handler in `userprog/exception.c`
 - Only check that a user pointer is below PHYS_BASE
 - Invalid pointers will trigger a page fault
 - See 3.1.5 [Accessing User Memory] for more details

Task 3: System Call Infrastructure

- Implement `syscall_handler()` in `syscall.c`
- What does this involve?
 - Read syscall number at stack pointer
 - SP is accessible as “esp” member of the struct `intr_frame *f` passed to `syscall_handler()`
 - Read some number of arguments above the stack pointer, depending on the syscall
 - Dispatch to the desired function
 - Return the value to the user in `f->eax`

Hints

- Syscall numbers are defined in `lib/syscall-nr.h`
 - You will not be implementing all the calls -- see 3.3.4 [System Calls] for a list of required calls for Project 2
- Many of the syscalls involve file system functionality
 - A simple filesys implementation is provided (you do NOT need to modify this, but familiarize yourself with the interfaces in `filesys.h` and `file.h`)
 - The file system is not thread-safe
 - Use coarse synchronization to ensure that any file system code is a critical section
 - Syscalls take “file descriptors” as arguments, but the Pintos file system uses `struct file *`s
 - You must design a proper mapping scheme

System Calls and Filesys (cont.)

- Reading from the keyboard and writing to the console are special cases with special file descriptors
- “write” syscall with fd `STDOUT_FILENO`
 - Use `putbuf(...)` or `putchar(...)` in `lib/kernel/console.c`
- “read” syscall with fd `STDIN_FILENO`
 - Use `input_getc(...)` in `devices/input.h`

Task 4: Implement System Calls (1)

■ System calls related to processes

```
void    exit (int status);  
pid_t   exec (const char *cmd_line);  
int     wait (pid_t pid);
```

- All of a process's resources must be freed on exit()
- The child can exit() before the parent performs wait()
- A process can perform wait() only for its children
- Wait() can be called twice for the same process
 - The second wait() should fail
- Nested waits are possible: $A \rightarrow B, B \rightarrow C$
- Pintos should not terminate until the initial process exits

==> Start with exit()

Hints

- `int wait (pid_t pid)`
 - caller blocks until the child process corresponding to `pid` exits
 - use synchronization primitives rather than `thread_block()`
 - returns the exit status of the child or -1 in certain cases
 - if `wait` has already been successfully called on the child
 - if `pid` does not reference a child of the caller
 - many cases to think about: multiple waits, nested waits, etc.
 - suggestion: implement `process_wait()`, and then `wait()` on top of `process_wait()`
 - involves the most work of all the syscalls
- `void exit (int status)`
 - terminate user program, return status to the kernel
 - print a termination message
 - if a child is exiting, communicate exit status back to the parent who called `wait`

Hints

- `pid_t exec (const char *cmd_line)`
 - behaves like Unix's `fork()` + `execve(...)`
 - creates a child process running the program in `cmd_line`
 - must not return until the new process has successfully loaded or failed (use synchronization to ensure this)

Task 4: Implement System Calls (2)

■ System calls related to files

```
bool    create (const char *file, unsigned initial_size);
bool    remove (const char *file);
int     open  (const char *file);
int     filesize (int fd);
int     read  (int fd, void *buffer, unsigned size);
int     write (int fd, void *buffer, unsigned size);
void    seek  (int fd, unsigned position);
unsigned tell  (int fd);
void    close (int fd);
```

- create()/remove()/open() work on file names
- The rest of them work on file descriptors

Hints

- **Implementing system calls related to files**
 - No need to change the code in the `filesys` directory
 - The existing routines in the `filesys` directory work on the "file" structure (`struct file *`)
 - Maintain a mapping structure from a file descriptor to the corresponding "file" structure
 - Deny writes to a running process's executable file
 - Ensure only one process at a time is executing the file system code

Task 5: Denying Writes to Executables

- Pintos should not allow code that is currently running to be modified
 - Use `file_deny_write()` to prevent writes to an open file
 - Note: closing a file will re-enable writes, so an executable file must be kept open as long as the process is running

Suggested Initial Strategy

- Make a disk and add a few programs (like args-*)
- Temporarily set up stack to avoid immediate page faults
 - In `setup_stack()` of `process.c`, change `*esp = PHYS_BASE` to `*esp = PHYS_BASE - 12`
 - This will allow execution of programs that take no args
- Implement safe user memory access
- Set up basic syscall infrastructure
 - Read the syscall number, dispatch to function
- Implement **exit** syscall
- Implement **write** syscall to `STDOUT_FILENO`
 - No tests will pass until you can functionally write to the console
- Change `process_wait()` to an infinite loop
 - The stub implementation exits immediately, so Pintos will power off before any processes can run

Creating Disk & Execute Programs

- In order to run programs, you will need to create simulated disks containing user code
- How? From userprog/build folder:
 - `pintos-mkdisk filesys.dsk --fileysys-size=2` /* creates a 2 MB disk */
 - `pintos -f -q` /* formats the disk */
 - `pintos -p ../examples/echo -a echo -- -q` /* copies ../examples/echo to disk and names it "echo" */
 - `pintos -q run 'echo x'` /* runs the program with arg x */
- Example programs are in `src/examples`, test programs are in `src/userprog/build/tests`
 - Run "make" in these directories to build the programs
- You may want to make backup disks, in case the disk gets trashed (there is no fileysys recovery tool)

Workspace

```
threads/thread.c      |    13
threads/thread.h      |    26 +
userprog/exception.c  |     8
userprog/process.c    |   247 ++++++++--
userprog/syscall.c    |   468 ++++++++--
userprog/syscall.h    |     1
6 files changed, 725 insertions(+), 38 deletions(-)
```

Step 4: Testing

Pintos provides a very systematic testing suite for your project:

1. Run all tests:

```
$ make check
```

2. Run individual tests:

```
$ make tests/userprog/args-none.result
```

3. Run the grading script:

```
$ make grade
```

Sample Functionality Test

```
/* This program echoes its command-line arguments */
```

```
int  
main (int argc, char *argv[])  
{  
    int i;  
    msg ("begin");  
    msg ("argc = %d", argc);  
    for (i = 0; i <= argc; i++)  
        if (argv[i] != NULL)  
            msg ("argv[%d] = '%s'", i, argv[i]);  
        else  
            msg ("argv[%d] = null", i);  
    msg ("end");  
    return 0;  
}
```

Expected output for 'args 1 2'

```
begin  
argc=3  
argv[0] = 'args'  
argv[1] = '1'  
argv[2] = '2'  
argv[3] = null  
end
```

Sample Robustness Test

```
/* This program attempts to read memory at an address that is not mapped.  
   This should terminate the process with a -1 exit code. */
```

```
#include "tests/lib.h"  
#include "tests/main.h"
```

```
void
```

```
test_main (void)
```

```
{
```

```
    msg ("Congratulations - you have successfully dereferenced NULL: %d",  
         *(int *)NULL);
```

```
    fail ("should have exited with -1");
```

```
}
```

Expected output:
bad-read: exit(-1)

Using GDB

- You can use GDB to debug user code
- Start GDB as usual, then do
`(gdb) loadusersymbols <userprog name>`
- You can set breakpoints and inspect data as usual
- Note: a name that appears in both the kernel and the user program will refer to the kernel
 - Fix this by typing `pintos-gdb <userprog name>`
 - Then, `loadusersymbols kernel.o`

Grading

10% of your project grade will be from the **design document**, and 90% from the successful **implementation** of the following components:

- Functionality of system calls, passes all twenty-eight (28) tests.
- Robustness of system calls, passes all thirty-four (34) tests.
- Functionality of features that VM might break, passes one (1) test.
- Functionality of base file system, passes all thirteen (13) tests.

You can use “**make grade**” to test your implementation and to see what grade you would get from the implementation component.

Step 5: Design Document

Use the template in Section 3.3.1 of the Pintos documentation.

```
| PROJECT 2: USER PROGRAMS |  
| DESIGN DOCUMENT         |  
+-----+  
  
---- GROUP ----  
  
>> Fill in the names and email addresses of your group members.  
  
FirstName LastName <email@domain.example>  
FirstName LastName <email@domain.example>  
FirstName LastName <email@domain.example>  
  
---- PRELIMINARIES ----  
  
>> If you have any preliminary comments on your submission, notes for the  
>> TAs, or extra credit, please give them here.  
  
>> Please cite any offline or online sources you consulted while  
>> preparing your submission, other than the Pintos documentation, course  
>> text, lecture notes, and course staff.  
  
ARGUMENT PASSING  
=====
```

```
---- DATA STRUCTURES ----  
  
>> A1: Copy here the declaration of each new or changed 'struct' or  
>> 'struct' member, global or static variable, 'typedef', or  
>> enumeration. Identify the purpose of each in 25 words or less.  
  
---- ALGORITHMS ----  
  
>> A2: Briefly describe how you implemented argument parsing. How do  
>> you arrange for the elements of argv[] to be in the right order?  
>> How do you avoid overflowing the stack page?
```

Submission

1. Design document

- due by **April 22nd @11:59PM**
- submit using your **submit_cse421** or **submit_cse521** script

2. All source code (the full source tree)

- due by **May 13th @11:59PM**
- submit via **AutoLab** auto-grading system

Late Policy

Up to 24 hour late submission is accepted with 20% penalty:

❖ before 11:59pm on May 13th -> no penalty

between 11:59pm on May 13th and 11:59pm on May 14th -->
20% penalty

after 11:59pm on May 14th --> submission not accepted!

Acknowledgements

- Pintos Manual
- Pintos Notes and Slides by A. He (Stanford), A. Romano (Stanford), J. Sundararaman & X. Liu (Virginia Tech), J. Kim (SKKU)